


Ethical scenarios drive GPT-4o's inconsistency

Mean NCI averaged over 10 scenarios × 4 temperatures per cell (3 temperatures for Claude).



Kruskal-Wallis on scenario-level means: $H(2) = 3.61$, $p = .16$ (omnibus n.s.). Pattern is descriptive and model-specific.