


Consistency and accuracy are not the same — the relationship is model-specific

Each dot = one (scenario × temperature) cell. Hollow circles = NCI = 1.0 with accuracy = 0 ('reliably wrong').

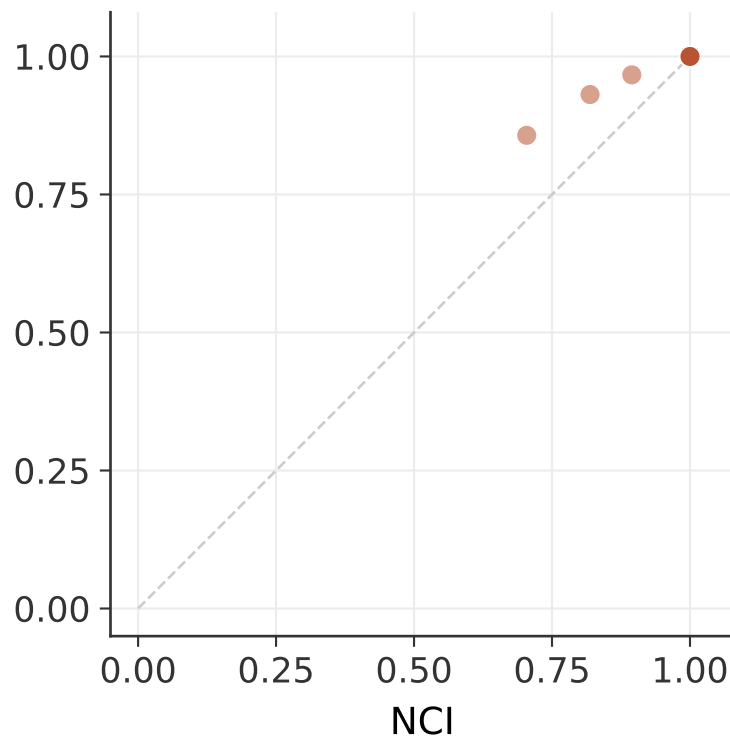
GPT-4o

Pearson $r = 0.50$ $p < .001$



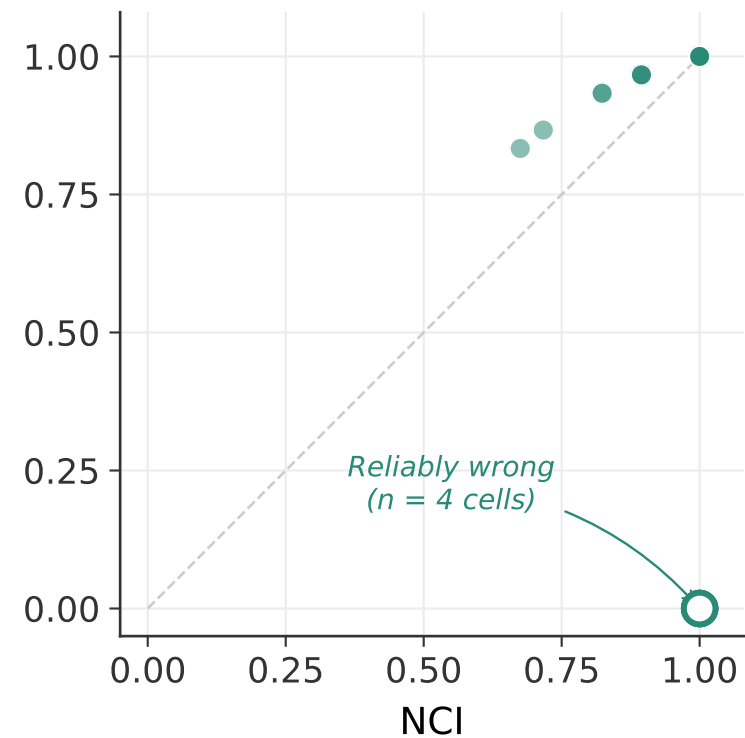
Claude Opus 4.5

Pearson $r = 0.99$ $p < .001$



Gemini 2.5 Flash-Lite

Pearson $r = 0.07$ $p = .42$



Dashed line: NCI = accuracy reference. GPT-4o $n = 120$ cells; Claude $n = 90$ ($T = 1.5$ unavailable); Gemini $n = 120$.